Responsiveness under Working Conditions
draft-ietf-ippm-responsiveness-01

Abstract

   For many years, a lack of responsiveness, variously called lag,
   latency, or bufferbloat, has been recognized as an unfortunate, but
   common, symptom in today's networks. Even after a decade of work on
   standardizing technical solutions, it remains a common problem for
   the end users.

   Everyone "knows" that it is "normal" for a video conference to have
   problems when somebody else at home is watching a 4K movie or
   uploading photos from their phone. However, there is no technical
   reason for this to be the case. In fact, various queue management
   solutions (fq_codel, cake, PIE) have solved the problem.

   Our networks remain unresponsive, not from a lack of technical
   solutions, but rather a lack of awareness of the problem and its
   solutions. We believe that creating a tool whose measurement matches
   people's everyday experience will create the necessary awareness, and
   result in a demand for products that solve the problem.

   This document specifies the "RPM Test" for measuring responsiveness.
   It uses common protocols and mechanisms to measure user experience
   specifically when the network is under working conditions. The
   measurement is expressed as "Round-trips Per Minute" (RPM) and should
   be included with throughput (up and down) and idle latency as
   critical indicators of network quality.

Status of This Memo

Table of Contents

1. Introduction

For many years, a lack of responsiveness, variously called lag,
latency, or bufferbloat, has been recognized as an unfortunate, but
common, symptom in today's networks [Bufferbloat]. Solutions like
fq_codel [RFC8290] or PIE [RFC8033] have been standardized and are to
some extent widely implemented. Nevertheless, people still suffer
from bufferbloat.

Although significant, the impact on user experience can be transitory
-- that is, its effect is not always visible to the user. Whenever a
network is actively being used at its full capacity, buffers can fill
up and create latency for traffic. The duration of those full
buffers may be brief: a medium-sized file transfer, like an email
attachment or uploading photos, can create bursts of latency spikes.
An example of this is lag occurring during a videoconference, where a
connection is briefly shown as unstable.

These short-lived disruptions make it hard to narrow down the cause.
We believe that it is necessary to create a standardized way to
measure and express responsiveness.

Existing network measurement tools could incorporate a responsiveness
measurement into their set of metrics. Doing so would also raise the
awareness of the problem and would help establish a new expectation
that the standard measures of network quality should -- in addition
to throughput and idle latency -- also include latency under load,
or, as we prefer to call it, responsiveness under working conditions.

1.1. Terminology

A word about the term "bufferbloat" -- the undesirable latency that
comes from a router or other network equipment buffering too much
data. This document uses the term as a general description of bad
latency, using more precise wording where warranted.

"Latency" is a poor measure of responsiveness, since it can be hard
for the general public to understand. The units are unfamiliar
("what is a millisecond?") and counterintuitive ("100 msec -- that
sounds good -- it's only a tenth of a second!").

Instead, we create the term "Responsiveness under working conditions" to make it clear that we are measuring all, not just idle, conditions, and use "round-trips per minute" as the metric. The advantage of round-trips per minute are two-fold: First, it allows for a metric that is "the higher the better". This kind of metric is often more intuitive for end-users. Second, the range of the values tends to be around the 4-digit integer range which is also a value easy to compare and read, again allowing for a more intuitive use. Finally, we abbreviate the measurement to "RPM", a wink to the "revolutions per minute" that we use for car engines.

This document defines an algorithm for the "RPM Test" that explicitly measures responsiveness under working conditions.

2. Design Constraints

There are many challenges around measurements on the Internet. They include the dynamic nature of the Internet, the diverse nature of the traffic, the large number of devices that affect traffic, and the difficulty of attaining appropriate measurement conditions.

Internet paths are changing all the time. Daily fluctuations in the demand make the bottlenecks ebb and flow. To minimize the variability of routing changes, it's best to keep the test duration relatively short.

TCP and UDP traffic, or traffic on ports 80 and 443, may take significantly different paths on the Internet and be subject to entirely different Quality of Service (QoS) treatment. A good test will use standard transport-layer traffic -- typical for people's use of the network -- that is subject to the transport's congestion control that might reduce the traffic's rate and thus its buffering in the network.

Traditionally, one thinks of bufferbloat happening on the routers and switches of the Internet. However, the networking stacks of the clients and servers can have huge buffers. Data sitting in TCP sockets or waiting for the application to send or read causes artificial latency, and affects user experience the same way as "traditional" bufferbloat.

Finally, it is crucial to recognize that significant queueing only happens on entry to the lowest-capacity (or "bottleneck") hop on a network path. For any flow of data between two communicating devices, there is always one hop along the path where the capacity available to that flow at that hop is the lowest among all the hops of that flow's path at that moment in time. It is important to understand that the existence of a lowest-capacity hop on a network

path is not itself a problem. In a heterogeneous network like the
Internet it is inevitable that there must necessarily be some hop
along the path with the lowest capacity for that path. If that hop
were to be improved to make it no longer the lowest-capacity hop,
then some other hop would become the new lowest-capacity hop for that
path. In this context a "bottleneck" should not be seen as a problem
to be fixed, because any attempt to "fix" the bottleneck is futile --
such a "fix" can never remove the existence of a bottleneck on a
path; it just moves the bottleneck somewhere else. Arguably, this
heterogeneity of the Internet is one of its greatest strengths.
Allowing individual technologies to evolve and improve at their own
pace, without requiring the entire Internet to change in lock-step,
has enabled enormous improvements over the years in technologies like
DSL, cable modems, Ethernet, and Wi-Fi, each advancing independently
as new developments became ready. As a result of this flexibility we
have moved incrementally, one step at a time, from 56kb/s dial-up
modems in the 1990s to Gb/s home Internet service and Gb/s wireless
connectivity today.

Note that in a shared datagram network, conditions do not remain
static. The hop that is the current bottleneck may change from
moment to moment. For example, changes in other traffic may result
in changes to a flow's share of a given hop. A user moving around
may cause the Wi-Fi transmission rate to vary widely, from a few Mb/s
when far from the Access Point, all the way up to Gb/s or more when
close to the Access Point.

Consequently, if we wish to enjoy the benefits of the Internet's
great flexibility, we need software that embraces and celebrates this
diversity and adapts intelligently to the varying conditions it
encounters.

Because significant queueing only happens on entry to the bottleneck
hop, the queue management at this critical hop of the path almost
entirely determines the responsiveness of the entire flow. If the
bottleneck hop's queue management algorithm allows an excessively
large queue to form, this results in excessively large delays for
packets sitting in that queue awaiting transmission, significantly
degrading overall user experience.

In order to discover the depth of the buffer at the bottleneck hop,
the RPM Test mimics normal network operations and data transfers, to
cause this bottleneck buffer to fill to capacity, and then measures
the resulting end-to-end latency under these operating conditions. A
well managed bottleneck queue keeps its queue occupancy under
control, resulting in consistently low round-trip time and
consistently good responsiveness. A poorly managed bottleneck queue
will not.

3. Goals

   The algorithm described here defines an RPM Test that serves as a
   good proxy for user experience. This means:

   1. Today's Internet traffic primarily uses HTTP/2 over TLS. Thus,
       the algorithm should use that protocol.

       As a side note: other types of traffic are gaining in popularity
       (HTTP/3) and/or are already being used widely (RTP). Traffic
       prioritization and QoS rules on the Internet may subject traffic
       to completely different paths: these could also be measured
       separately.

   2. The Internet is marked by the deployment of countless middleboxes
       like transparent TCP proxies or traffic prioritization for
       certain types of traffic. The RPM Test must take into account
       their effect on TCP-handshake [RFC0793], TLS-handshake, and
       request/response.

   3. The test result should be expressed in an intuitive, nontechnical
       form.

   4. Finally, to be useful to a wide audience, the measurement should
       finish within a short time frame. Our target is 20 seconds.

4. Measuring Responsiveness Under Working Conditions

   To make an accurate measurement, the algorithm must reliably put the
   network in a state that represents those "working conditions".
   During this process, the algorithm measures the responsiveness of the
   network. The following explains how the former and the latter are
   achieved.

4.1. Working Conditions

   There are many different ways to define the state of "working
   conditions" to measure responsiveness. There is no one true answer
   to this question. It is a tradeoff between using realistic traffic
   patterns and pushing the network to its limits.

   The working conditions we try to achieve is a scenario where the path
   between the measuring endpoints is utilized at its full end-to-end
   capacity. An ideal sender could send at just this link-speed without
   building a queue on the bottleneck. Thus, in order to measure the
   worst-case responsiveness we need to ensure that a queue is building
   up on the bottleneck, meaning that responsiveness is at its worst.

In this document we aim to generate a realistic traffic pattern by using standard HTTP transactions but exploring the worst-case scenario by creating multiple of these transactions and using very large data objects in these HTTP transactions.

This allows to create a stable state of working conditions during which the bottleneck of the path between client and server has its buffer filled up entirely, without generating DoS-like traffic patterns (e.g., intentional UDP flooding). This creates a realistic traffic mix representative of what a typical user's network experiences in normal operation.

Finally, as end-user usage of the network evolves to newer protocols and congestion control algorithms, it is important that the working conditions also can evolve to continuously represent a realistic traffic pattern.

4.1.1. From single-flow to multi-flow

A single TCP connection may not be sufficient to reach the capacity and full buffer occupancy of a path quickly. Using a 4MB receive window, over a network with a 32 ms round-trip time, a single TCP connection can achieve up to 1Gb/s throughput. Additionally, deep buffers along the path between the two endpoints may be significantly larger than 4MB. TCP allows larger receive window sizes, up to 1GB. However, most transport stacks aggressively limit the size of the receive window to avoid consuming too much memory.

Thus, the only way to achieve full capacity and full buffer occupancy on those networks is by creating multiple connections, allowing to actively fill the bottleneck's buffer to achieve maximum working conditions.

Even if a single TCP connection would be able to fill the bottleneck's buffer, it may take some time for a single TCP connection to ramp up to full speed. One of the goals of the RPM test is to quickly load the network, take its measurements, and then finish. Finally, traditional loss-based TCP congestion control algorithms react aggressively to packet loss by reducing the congestion window. This reaction (intended by the protocol design) decreases the queueing within the network, making it harder to determine the depth of the bottleneck queue reliably.

The purpose of the RPM Test is not to productively move data across the network in a useful way, the way a normal application does. The purpose of the RPM Test is, as quickly as possible, to simulate a representative traffic load as if real applications were doing sustained data transfers, measure the resulting round-trip time

occurring under those realistic conditions, and then end the test.
Because of this, using multiple simultaneous parallel connections
allows the RPM test to complete its task more quickly, in a way that
overall is less disruptive and less wasteful of network capacity than
a test using a single TCP connection that would take longer to bring
the bottleneck hop to a stable saturated state.

4.1.2. Parallel vs Sequential Uplink and Downlink

Poor responsiveness can be caused by queues in either (or both) the
upstream and the downstream direction. Furthermore, both paths may
differ significantly due to access link conditions (e.g., 5G
downstream and LTE upstream) or the routing changes within the ISPs.
To measure responsiveness under working conditions, the algorithm
must explore both directions.

One approach could be to measure responsiveness in the uplink and
downlink in parallel. It would allow for a shorter test run-time.

However, a number of caveats come with measuring in parallel:

* Half-duplex links may not permit simultaneous uplink and downlink
   traffic. This means the test might not reach the path's capacity
   in both directions at once and thus not expose all the potential
   sources of low responsiveness.

* Debuggability of the results becomes harder: During parallel
   measurement it is impossible to differentiate whether the observed
   latency happens in the uplink or the downlink direction.

Thus, we recommend testing uplink and downlink sequentially.
Parallel testing is considered a future extension.

4.1.3. Reaching full buffer utilization

The RPM Test gradually increases the number of TCP connections and
measures "goodput" (the sum of actual data transferred across all
connections in a unit of time) as well as responsiveness
continuously. When both goodput and responsiveness stop changing, it
means that the test managed to fill the buffer of the bottleneck. At
this point we are creating the worst-case scenario within the limits
of the realistic traffic pattern.

The algorithm notes that throughput increases rapidly until TCP
connections complete their TCP slow-start phase. At that point,
throughput eventually stalls, often due to receive window
limitations, particularly in cases of high network bandwidth, high
network round-trip time, low receive window size, or a combination of

all three. The only means to further increase throughput is by
adding more TCP connections to the pool of load-generating
connections. If new connections leave the throughput the same, full
link utilization has been reached and -- more importantly -- the
working condition is stable.

4.2. Measuring Responsiveness

Measuring responsiveness during the previously explained working
conditions creation is a continuous process during the duration of
the test. It requires a sufficiently large sample-size to have
confidence in the results.

The measurement of the responsiveness happens by sending probe-
requests for a small object. There are two types of probe requests:

1. A HTTP GET request on a separate connection ("foreign probes").
    This test mimics the time it takes for a web browser to connect
    to a new web server and request the first element of a web page
    (e.g., "index.html"), or the startup time for a video streaming
    client to launch and begin fetching media.

2. A HTTP GET request multiplexed on the load-generating connections
    ("self probes"). This test mimics the time it takes for a video
    streaming client to skip ahead to a different chapter in the same
    video stream, or for a navigation client to react and fetch new
    map tiles when the user scrolls the map to view a different area.
    In a well functioning system fetching new data over an existing
    connection should take less time than creating a brand new TLS
    connection from scratch to do the same thing.

Foreign probes will provide 3 sets of data-points. First, the
duration of the TCP-handshake (noted hereafter as tcp_foreign).
Second, the TLS round-trip-time (noted tls_foreign). For this, it is
important to note that different TLS versions have a different number
of round-trips. Thus, the TLS establishment time needs to be
normalized to the number of round-trips the TLS handshake takes until
the connection is ready to transmit data. And third, the HTTP
elapsed time between issuing the GET request for a 1-byte object and
receiving the entire response (noted http_foreign).

Self probes will provide a single data-point for the duration of time
between when the HTTP GET request for the 1-byte object is issued on
the load-generating connection and the full HTTP response has been
received (noted http_self).

tcp_foreign, tls_foreign, http_foreign and http_self are all measured
in milliseconds.

The more probes that are sent, the more data is available for
calculation. In order to generate as much data as possible, the
methodology requires a client to issue these probes every 100ms. For
the probes on the load-generating connections, the client needs to
use one of the initial load-generating connections. This means that
every 100ms, 2 probes are being evaluated. The total amount of data
used for these probes would be no more than about 50KB worth of data
within one second.

4.2.1. Aggregating the Measurements

The algorithm produces sets of 4 times for each probe, namely:
tcp_foreign, tls_foreign, http_foreign, http_self (from the previous
section). Each of these sets will have a large number of samples.
Use the following methodology to calculate a single RPM value from
these data:

1. Among each set, take the 90th percentile, thus resulting in 4
    individual numbers (tcp_foreign_p90, tls_foreign_p90,
    http_foreign_p90, http_self_p90).

2. Calculate the RPM as the weighted mean:

Responsiveness = 60000 /
((1/3*tcp_foreign_p90 + 1/3*tls_foreign_p90 + 1/3*http_foreign_p90 +
  http_self_p90)/2)

This responsiveness value presents round-trips per minute (RPM).

4.3. Final Algorithm

Considering the previous two sections, where we explain what the
meaning of working conditions is and the definition of
responsiveness, we can design the final algorithm. In order to
measure the worst-case latency we need to transmit traffic at the
full capacity of the path as well as ensure the buffers are filled to
the maximum. We can achieve this by continuously adding HTTP
sessions to the pool of connections in a 1-second interval. This
will ensure that we quickly reach capacity and full buffer occupancy.
We need to continuously measure goodput and responsiveness and as
soon as we detect stability for both metrics we can ensure that the
full working conditions have been reached.

The following algorithm reaches working conditions of a network by
using HTTP/2 upload (POST) or download (GET) requests of infinitely
large files. The algorithm is the same for upload and download and
uses the same term "load-generating connection" for each. The
actions of the algorithm take place at regular intervals. For the
current draft the interval is defined as one second.

Where

*   i: The index of the current interval. The variable i is
    initialized to 0 when the algorithm begins and increases by one
    for each interval.

*   instantaneous aggregate goodput at interval p: The number of total
    bytes of data transferred within interval p, divided by the
    interval duration. If p is negative (i.e., a time interval
    logically prior to the start of the test beginning, used in moving
    average calculations), the number of total bytes of data
    transferred within that interval is considered to be 0.

*   moving average aggregate goodput at interval p: The number of
    total bytes of data transferred within interval p and the three
    immediately preceding intervals, divided by four times the
    interval duration.

the steps of the algorithm are:

*   Create a load-generating connection.

*   Start probing for responsiveness every 100ms, as described in the
    previous section.

*   At each interval:

    -   Create an additional load-generating connection.

    -   Compute the instantaneous aggregate goodput at interval i.

    -   Compute the moving average aggregate goodput at interval i.

    -   Compute the responsiveness

    -   If the moving average aggregate goodput at interval i is more
        than a 5% increase over the moving average aggregate goodput at
        interval i - 1, the network has not yet reached full link
        utilization. Continue for 4 more iterations.

-   If the responsiveness at interval i is more than a 5% reduction
    over the responsiveness at interval i - 1, the network has not
    yet reached full buffer occupancy. Continue for 4 more
    iterations.

In Section 3, it is mentioned that one of the goals is that the test
finishes within 20 seconds. It is left to the implementation what to
do when stability is not reached within that time-frame. For
example, an implementation might gather a provisional responsiveness
measurement or let the test run for longer.

Finally, if at any point one of these connections terminates with an
error, the test should be aborted.

5. Interpreting responsiveness results

The described methodology uses a high-level approach to measure
responsiveness. By executing the test with regular HTTP requests a
number of elements come into play that will influence the result.
Contrary to more traditional measurement methods the responsiveness
metric is not only influenced by the properties of the network but
can significantly be influenced by the properties of the client and
the server implementations. This section describes how the different
elements influence responsiveness and how a user may differentiate
them when debugging a network.

5.1. Elements influencing responsiveness

Due to the HTTP-centric approach of the measurement methodology a
number of factors come into play that influence the results. Namely,
the client-side networking stack (from the top of the HTTP-layer all
the way down to the physical layer), the network (including potential
transparent HTTP "accelerators"), and the server-side networking
stack. The following outlines how each of these contributes to the
responsiveness.

5.1.1. Client side influence

As the driver of the measurement, the client-side networking stack
can have a large influence on the result. The biggest influence of
the client comes when measuring the responsiveness in the uplink
direction. Load-generation will cause queue-buildup in the transport
layer as well as the HTTP layer. Additionally, if the network's
bottleneck is on the first hop, queue-buildup will happen at the
layers below the transport stack (e.g., NIC firmware).

Each of these queue build-ups may cause latency and thus low responsiveness. A well designed networking stack would ensure that queue-buildup in the TCP layer is kept at a bare minimum with solutions like TCP_NOTSENT_LOWAT [draft-ietf-tcpm-rfc793bis]. At the HTTP/2 layer it is important that the load-generating data is not interfering with the latency-measuring probes. For example, the different streams should not be stacked one after the other but rather be allowed to be multiplexed for optimal latency. The queue-buildup at these layers would only influence latency on the probes that are sent on the load-generating connections.

Below the transport layer many places have a potential queue build-up. It is important to keep these queues at reasonable sizes or that they implement techniques like FQ-Codel. Depending on the techniques used at these layers, the queue build-up can influence latency on probes sent on load-generating connections as well as separate connections. If flow-queuing is used at these layers, the impact on separate connections will be negligible.

5.1.2. Network influence

The network obviously is a large driver for the responsiveness result. Propagation delay from the client to the server as well as queuing in the bottleneck node will cause latency. Beyond these traditional sources of latency, other factors may influence the results as well. Many networks deploy transparent TCP Proxies, firewalls doing deep packet-inspection, HTTP "accelerators",... As the methodology relies on the use of HTTP/2, the responsiveness metric will be influenced by such devices as well.

The network will influence both kinds of latency probes that the responsiveness tests sends out. Depending on the network's use of Smart Queue Management and whether this includes flow-queuing or not, the latency probes on the load-generating connections may be influenced differently than the probes on the separate connections.

5.1.3. Server side influence

Finally, the server-side introduces the same kind of influence on the responsiveness as the client-side, with the difference that the responsiveness will be impacted during the downlink load generation.

5.2. Root-causing Responsiveness

Once an RPM result has been generated one might be tempted to try to
localize the source of a potential low responsiveness. The
responsiveness measurement is however aimed at providing a quick,
top-level view of the responsiveness under working conditions the way
end-users experience it. Localizing the source of low responsiveness
involves however a set of different tools and methodologies.

Nevertheless, the responsiveness test allows to gain some insight
into what the source of the latency is. The previous section
described the elements that influence the responsiveness. From there
it became apparent that the latency measured on the load-generating
connections and the latency measured on separate connections may be
different due to the different elements.

For example, if the latency measured on separate connections is much
less than the latency measured on the load-generating connections, it
is possible to narrow down the source of the additional latency on
the load-generating connections. As long as the other elements of
the network don't do flow-queueing, the additional latency must come
from the queue build-up at the HTTP and TCP layer. This is because
all other bottlenecks in the network that may cause a queue build-up
will be affecting the load-generating connections as well as the
separate latency probing connections in the same way.

6. RPM Test Server API

The RPM measurement is built upon a foundation of standard protocols:
IP, TCP, TLS, HTTP/2. On top of this foundation, a minimal amount of
new "protocol" is defined, merely specifying the URLs that used for
GET and PUT in the process of executing the test.

Both the client and the server MUST support HTTP/2 over TLS. The
client MUST be able to send a GET request and a POST. The server
MUST be able to respond to both of these HTTP commands. The server
MUST have the ability to provide content upon a GET request. The
server MUST use a packet scheduling algorithm that minimizes internal
queueing to avoid affecting the client's measurement.

As clients and servers become deployed that use L4S congestion
control (e.g., TCP Prague with ECT(1) packet marking), for their
normal traffic when it is available, and fall back to traditional
loss-based congestion controls (e.g., Reno or CUBIC) otherwise, the
same strategy SHOULD be used for RPM test traffic. This is
RECOMMENDED so that the synthetic traffic generated by the RPM test
mimics real-world traffic for that server.

Delay-based congestion-control algorithms (e.g., Vegas, FAST, BBR)
SHOULD NOT be used for RPM test traffic because they take much longer
to discover the depth of the bottleneck buffers. Delay-based
congestion-control algorithms seek to mitigate the effects of
bufferbloat, by detecting and responding to early signs of increasing
round-trip delay, and reducing the amount of data they have in flight
before the bottleneck buffer fills up and overflows. In a world
where bufferbloat is common, this is a pragmatic mitigation to allow
software to work better in that environment. However, that approach
does not fix the underlying problem of bufferbloat; it merely avoids
it in some cases, and allows the problem in the network to persist.
For a diagnostic tool made to identify symptoms of bufferbloat in the
network so that they can be fixed, using a transport protocol
explicitly designed to mask those symptoms would be a poor choice,
and would require the test to run for much longer to deliver the same
results.

The server MUST respond to 4 URLs:

1. A "small" URL/response: The server must respond with a status
    code of 200 and 1 byte in the body. The actual message content
    is irrelevant. The server SHOULD specify the content-type as
    application/octet-stream. The server SHOULD minimize the size,
    in bytes, of the response fields that are encoded and sent on the
    wire.

2. A "large" URL/response: The server must respond with a status
    code of 200 and a body size of at least 8GB. The server SHOULD
    specify the content-type as application/octet-stream. The body
    can be bigger, and may need to grow as network speeds increases
    over time. The actual message content is irrelevant. The client
    will probably never completely download the object, but will
    instead close the connection after reaching working condition and
    making its measurements.

3. An "upload" URL/response: The server must handle a POST request
    with an arbitrary body size. The server should discard the
    payload. The actual POST message content is irrelevant. The
    client will probably never completely upload the object, but will
    instead close the connection after reaching working condition and
    making its measurements.

4. A configuration URL that returns a JSON [RFC8259] object with the
    information the client uses to run the test (sample below). The
    server SHOULD specify the content-type as application/json.
    Sample JSON:

```
{
  "version": 1,
  "urls": {
    "large_https_download_url":"https://nq.example.com/api/v1/large",
    "small_https_download_url":"https://nq.example.com/api/v1/small",
    "https_upload_url":        "https://nq.example.com/api/v1/upload"
  }
  "test_endpoint": "hostname123.provider.com"
}
```

All of the fields in the sample configuration are required except
"test_endpoint". If the test server provider can pin all of the
requests for a test run to a specific host in the service (for a
particular run), they can specify that host name in the
"test_endpoint" field.

The client begins the responsiveness measurement by querying for the
JSON configuration. This supplies the URLs for creating the load-
generating connections in the upstream and downstream direction as
well as the small object for the latency measurements.

7. RPM Test Server Discovery

It makes sense to host RPM Test Server instances in Internet Data
Centers where they can be accessed easily by users wishing to test
the quality of their Internet connection. However, when a user
performs an RPM test and determines that they are suffering from poor
RPM during download, the logical next question might be, "What's
causing my poor performance? Is it poor buffer management by my ISP?
Is it poor buffer management in my home Wi-Fi Access point?
Something else?"

To help an end user answer this question, it will be useful for home
gateway equipment to host RPM Test Server instances. In an example
configuration, a user may have cable modem service offering 100 Mb/s
download speed, connected via gigabit Ethernet to one or more Wi-Fi
access points in the home, which then offer service to Wi-Fi client
devices at different rates depending on distance, interference from
other traffic, etc. By having the cable modem itself host an RPM
Test Server instance, the user can then run a test between the cable
modem and their computer or smartphone, to help isolate whether
bufferbloat they are experiencing is occurring in equipment inside
the home (like their Wi-Fi access points) or somewhere outside the
home.

To aid in discoverability of these facilities, local RPM Test Server
instances SHOULD advertise the availability of service type [RFC6335]
"_nq._tcp" (Network Quality), via DNS-Based Service Discovery

[RFC6763], using Multicast DNS on its local link(s) [RFC6762]. Where
applicable, an RPM Test Server instance SHOULD also advertise the
availability of its service via unicast discovery, for discovery by
client devices not directly attached to the same link. Population of
the appropriate DNS zone with the relevant unicast discovery records
can be performed automatically using a Discovery Proxy [RFC8766], or
in some scenarios simply by having a human administrator manually
enter the required records. Similarly, a "cloud" service, providing
Internet hosting service for "example.com" could choose to include
the relevant DNS-SD records within the "example.com" domain [RFC6763]
to communicate to clients the list of available RPM Test Server
instances.

8. Security Considerations

   TBD

9. IANA Considerations

   IANA has been requested to record the service type "_nq._tcp"
   (Network Quality) for advertising and discovery of RPM Test Server
   instances.

10. Acknowledgments

   Special thanks go to Will Hawkins and Jeroen Schickendantz for their
   tireless enthousiasm around the project and their contributions to
   this I-D and the development of the Go responsiveness measurement
   tool. We would also like to thank Rich Brown for his editorial pass
   over this I-D. We also thank Erik Auerswald for his constructive
   feedback on the I-D.

11. Informative References

   [Bufferbloat]
             Gettys, J. and K. Nichols, "Bufferbloat: Dark Buffers in
             the Internet", Communications of the ACM, Volume 55,
             Number 1 (2012) , n.d..

   [draft-ietf-tcpm-rfc793bis]
             Eddy, W., "Transmission Control Protocol (TCP)
             Specification", Internet Engineering Task Force , n.d..

   [RFC0793] Postel, J., "Transmission Control Protocol", RFC 793,
             DOI 10.17487/RFC0793, September 1981,
             <https://www.rfc-editor.org/info/rfc793>.

   [RFC6335] Cotton, M., Eggert, L., Touch, J., Westerlund, M., and S.
              Cheshire, "Internet Assigned Numbers Authority (IANA)
              Procedures for the Management of the Service Name and
              Transport Protocol Port Number Registry", BCP 165,
              RFC 6335, DOI 10.17487/RFC6335, August 2011,
              <https://www.rfc-editor.org/info/rfc6335>.

   [RFC6762] Cheshire, S. and M. Krochmal, "Multicast DNS", RFC 6762,
              DOI 10.17487/RFC6762, February 2013,
              <https://www.rfc-editor.org/info/rfc6762>.

   [RFC6763] Cheshire, S. and M. Krochmal, "DNS-Based Service
              Discovery", RFC 6763, DOI 10.17487/RFC6763, February 2013,
              <https://www.rfc-editor.org/info/rfc6763>.

   [RFC8033] Pan, R., Natarajan, P., Baker, F., and G. White,
              "Proportional Integral Controller Enhanced (PIE): A
              Lightweight Control Scheme to Address the Bufferbloat
              Problem", RFC 8033, DOI 10.17487/RFC8033, February 2017,
              <https://www.rfc-editor.org/info/rfc8033>.

   [RFC8259] Bray, T., Ed., "The JavaScript Object Notation (JSON) Data
              Interchange Format", STD 90, RFC 8259,
              DOI 10.17487/RFC8259, December 2017,
              <https://www.rfc-editor.org/info/rfc8259>.

   [RFC8290] Hoeiland-Joergensen, T., McKenney, P., Taht, D., Gettys,
              J., and E. Dumazet, "The Flow Queue CoDel Packet Scheduler
              and Active Queue Management Algorithm", RFC 8290,
              DOI 10.17487/RFC8290, January 2018,
              <https://www.rfc-editor.org/info/rfc8290>.

   [RFC8446] Rescorla, E., "The Transport Layer Security (TLS) Protocol
              Version 1.3", RFC 8446, DOI 10.17487/RFC8446, August 2018,
              <https://www.rfc-editor.org/info/rfc8446>.

   [RFC8766] Cheshire, S., "Discovery Proxy for Multicast DNS-Based
              Service Discovery", RFC 8766, DOI 10.17487/RFC8766, June
              2020, <https://www.rfc-editor.org/info/rfc8766>.

Appendix A. Example Server Configuration

   This section shows fragments of sample server configurations to host
   an responsiveness measurement endpoint.

A.1. Apache Traffic Server

   Apache Traffic Server starting at version 9.1.0 supports
   configuration as a responsiveness server. It requires the generator
   and the statichit plugin.

   The sample remap configuration file then is:

   map https://nq.example.com/api/v1/config \
       http://localhost/ \
       @plugin=statichit.so \
       @pparam=--file-path=config.example.com.json \
       @pparam=--mime-type=application/json

   map https://nq.example.com/api/v1/large \
       http://localhost/cache/8589934592/ \
       @plugin=generator.so

   map https://nq.example.com/api/v1/small \
       http://localhost/cache/1/ \
       @plugin=generator.so

   map https://nq.example.com/api/v1/upload \
       http://localhost/ \
       @plugin=generator.so

Authors' Addresses

   Christoph Paasch
   Apple Inc.
   One Apple Park Way
   Cupertino, California 95014,
   United States of America
   Email: cpaasch@apple.com

   Randall Meyer
   Apple Inc.
   One Apple Park Way
   Cupertino, California 95014,
   United States of America
   Email: rrm@apple.com

Stuart Cheshire
Apple Inc.
One Apple Park Way
Cupertino, California 95014,
United States of America
Email: cheshire@apple.com

Omer Shapira
Apple Inc.
One Apple Park Way
Cupertino, California 95014,
United States of America
Email: oesh@apple.com

Matt Mathis
Google, Inc
1600 Amphitheatre Parkway
Mountain View, CA 94043,
United States of America
Email: mattmathis@google.com